

# Scanner

quarta-feira, 14 de maio de 2025 18:55

- Crafting interpreters
  - Take in the source code
    - as a series of characters
    - group in into a series of chunks called tokens
  - Switch with a delusions of grandeur
  - Define how to run the code
    - Receive a source file and process it
  - Good first stage is to read the tokens and print it.
  - Not error handling but save file some information to find the line latter on
    - linefile
    - file name
  - Separate the code
    - The ones that identify the error and the one that reports them
  - Lexical analysis
    - scan through the list of characters and group them together into the smallest sequence that still represent something
    - Each blob of characters is called a lexeme
    - when we group the lexeme and bundle it together with the other data, the result is a token
    - Token
      - Keywords are part of the shape of a language grammar
        - ◆ We would like to know if the next word is a while for example
      - Parser could categorize tokens from the raw lexeme by comparing the strings ( slow and ugly )
      - At the point we recognize a lexeme, we also remember which kind of lexeme it represents ( we have a different type for each keyword, operator, bit, type ... )
        - ◆ Defined in an enum
      - Literal Value ( numbers and string )
      - Location information
        - ◆ Define a class token with all the information we need ( linefile for example. Type, lexeme, literal and line)
          - ◇ Also good to make a toString() method
      - Regular language and expressions
        - ◆ The core of a scanner is a loop
        - ◆ Starting at the first character of the source code
          - ◇ figures it out which lexeme it belongs
          - ◇ consume it and any following characters that are part of the lexeme
          - ◇ When reach the end of the lexeme, emits a token
          - ◇ Keeps doing that, eating characters and occasionally excreting tokens until reach the end of the source code
        - ◆ May consider to define a regexp for each kind of lexeme
          - ◇ The rules that determine how a particular language groups characters into lexeme are called lexical grammar
          - ◇ Tools like LEX are designed expressly to let you throw a handful of regexps and they give a complete scanner back ( won't delegating it ) but we can....
    - Scanner class
      - Store the raw source code into a string
      - define a list to be ready to fill with tokens ( not yet generated )
      - Loop each character in the string source code
      - When it run out of characters, append one final "end of file" token
        - ◆ not needed but makes it cleaner
        - ◆ Depends on a couple of fields,
          - ◇ start -> points to the first character of the lexeme
          - ◇ current -> points to the character currently been considered
            - ▶ offsets that index into the string to be looped
          - ◇ Line -> tracks what source line current is on ( can improve this... linefile)
            - ▶ Produce tokens that know their location
          - ◇ Helper to see if we reached the end
            - ▶ current >= source.length()
      - Recognizing lexemes
        - heart of the scanner
        - each turn of the loop we scan a single token
        - Starting with the single characters
          - ◆ consume the next character and pick a token type for it
            - ◇ Why consume the next?
              - ▶ don't think we need... maybe java starts list position at 1 and not 0
              - ▶ maybe we need... yeah... we need to get the next token... but this may change
              - ▶ ok.. we do need because this is the move forward of the while
          - ◆ add token, substring from the start to the current character
          - ◆ add the token to the list
      - Lexical errors
        - If not handled, will be silently discarded
          - ◆ Probably what i'm going to do
        - handled as a default case, or not supported
      - Operators

- Some of them requires 2 characters
        - ◆ In this case we match the first one and compare the next one... is it < followed by a =? then BANG\_EQUAL, otherwise just BANG
  - Longer lexemes
    - Comments \\\
      - ◆ When we find it... jump to the next line... **IMPORTANT**
      - ◆ Comments usually goes to the end of the line... so if finding it... create a comment token until the end of the line
      - ◆ I'll need to implement this
    - This is our general strategy for handling longer lexemes. After we detect the beginning of one, we shunt over to some lexeme-specific code that keeps eating characters until it sees the end.
    - We can't peek one or two ahead to match a lexeme... usually one is faster
    - Comments aren't meaningful lexemes and the parser doesn't want to deal with them ( don't call add Token() for comments )
      - ◆ When start looking for the next lexeme, start member get's reset and the comment's lexeme disappears
      - ◆ skip over other meaningless characters ( new lines and white spaces ) ( ' ' ) ( \r \n )
        - ◇ Just remember to increment a new line when skipping lines
        - ◇ For white spaces, simply go back to the beginning of the scan loop to start a new lexeme
  - String literals
    - Always begins with a specific character "
      - ◆ create a case for that
      - ◆ If doesn't find a matching end... generate an error
      - ◆ Consume characters until we hit the final "
      - ◆ Consider multiple line strings
      - ◆ When creating the token, we also produce the string value, that will be used later by the interpreter
  - Number literals
    - A series of digits optionally followed by a "." and more digits
    - To recognize a number, we look for any digit > "0" and < "9"
    - Separate method to consume the rest of the literal in case of number
    - consume as many digits as we find
      - ◆ For the integer and if a "." for the float part as well
      - ◆ Here we need a second lookahead, because we don't want to consume a "." if no digit after that
    - convert the lexeme to the numeric value
  - Reserved words and identifiers
    - we can't use a peek one ahead because the user can type a lot of things... like orNot and this is not an or logical operator
    - Use the maximal munch principle
      - ◆ when two lexical grammar rules can both match a chunk of code that the scanner is looking at, whichever one matches the most characters wins
      - ◆ it means we can't easily detect a reserved word until we've reached the end of what might instead be an identifier.
      - ◆ Reserved words are identifiers claimed by the language
    - we assume any lexeme starting with a letter or underscore is an identifier
      - ◆ if digit, number()
      - ◆ else if alpha() identifier()
        - ◇ isAlpha >='a' <='z' || >='A' && <='Z' || == '\_'
      - ◆ else unexpected character error
    - To handle keywords, we see if the identifier's lexeme is one of the reserved words
      - ◆ In this case, reserved word class
      - ◆ Define a set of reserved keywords for that in a map
      - ◆ so, after we scan an identifier, we check to see if it matches anything in the map
      - ◆ If matches, use the reserved keyword as token, otherwise it's just a user-defined identifier
- Type some expected and unexpected code and see if it produces the expected tokens

## • Engineering a compiler

- Scanner is the only pass in the compiler to touch every character of the input program
- Needs to be highly efficient
- First stage of a 3 part process that compilers use to understand the input program
- Scanner or lexical analyzer
  - reads a stream of characters and produces a stream of words
  - if the word is valid, the compiler assigns a syntactic category
- There are some tools for that but implementing it is simple and can be fast and robust
  - Most textbooks and courses use the generated scanners
  - most commercial compilers and open-source compilers use hand-crafted scanners
    - which can be faster because the implementation can optimize it as needed
- 3 different ways
  - table-driven scanner
  - direct-coded scanner
  - hand-coded approach
- microsyntax
  - specifies how to group characters into words and how to separate words that run together
  - Some times keywords and reserved words match the rule for an identifier but have special meanings
  - To recognize keywords, the scanner can either use dictionary lookup or encode the keywords directly into its microsyntax rule

- The compiler writer starts from a specification of the language microsyntax
- Implemented to requires just  $O(1)$  time per character
  - So they run in time proportional to the number of characters in the input stream
- Recognizing words
  - Let's identify the word new
    - Assuming a routine NextChar -> that returns the next character
  - The code test for n followed by e followed by w
  - Th code fragmentperforms one test per character